

Program 3 Spring 2018
20 Points
Rotation Puzzle Revisited AVL Trees

PART 1

You have been given the AVL tree code from your author as a starting point. Reading code is a valuable experience. On the job, you are often asked to modify/maintain existing code. You can't start over and do it your way. You must incorporate your changes into the existing code. **You are expected to understand the code that has been given to you.** Make the following changes:

1. Change all variable names to be meaningful names.
2. Make any changes needed to conform to our style guidelines.
3. Write a toString function which creates an indented version of the tree (as in program 2).
4. Implement the removeMin function. Remove the smallest node in the tree and rebalance. **This is to be your own work, not copied from ANY other source.**

The data to be stored in the AVL tree is to be generic and work for either an integer or a GameState (of the rotation puzzle variety).

Illustrate that the AVL tree is working properly by printing the tree (using toString) after **each of the following** groups of operations:

- Add: 1 3 5 7 9 11 2 4 8 (now print tree)
- Remove 7 (now print tree)
- Add 50 30 15 18 (now print tree)
- Remove Min (now print tree)
- Remove Min (now print tree)
- Remove Min (now print tree)
- Add 17 (now print tree)

PART 2

Use the AVL tree as a priority queue to improve your solution to the Rotation puzzle problem (of program 1).

A priority queue is a queue such that insertions are made in any order, but when you remove a node from the queue, it is the one with the best priority. Our priority will be an estimation of the total amount of work needed to solve the problem – so a lower score is preferred. Thus, we will first consider boards that “look better” to us.

In our previous solution, (program 1), we considered all one-move solutions before considering all two move solutions, and so on. If we are smarter in which solutions we consider, we can improve the functioning of the solution. **In this assignment, you will compare your previous solution to this new technique.**

We define a *state* of the game to contain at least:

- board

- cost so far: number of moves taken from initial state to reach current board
- expected cost to reach a solution
- String of moves indicating how the current board was achieved.

Best-first search. We describe a solution to the problem that illustrates a general artificial intelligence methodology known as the [A* search algorithm](#).

First, insert the initial state into a priority queue. Then, delete from the priority queue the state with the smallest expected cost, and insert onto the priority queue all neighboring states (those that can be reached in one more move) using an appropriate estimated cost for each. Repeat this procedure until the state removed from the priority queue is the goal state. The success of this approach hinges on the choice of estimated cost or *priority function*.

You can compute your “expected work function” any way you want as long it is reasonable and underestimates the real cost. Be creative.

Here is one approach. You can use any method you want. This is just one possibility. For each number, give each entry a zero (if it is in the right location), a 1 if it is in the correct row, but the wrong column, a 1 if it is in the correct column but the wrong row, and a 2 if it is in the wrong row and wrong column.

```
5 8 6
4 3 1
9 7 2
```

So for example, for the board above, the distances each cell is from the goal location is shown in red. Note, 5 has a distance of two because it is in the wrong row and wrong column.

5	8	6		2	1	1
4	3	1		0	2	2
9	7	2		1	1	2

If I sum these distances, I get 12. However, since one rotation can affect the distance of three different cells, we divide the total distance by 3 so that we have an underestimate of the actual number of moves required. (A* search does not work if you don’t use an underestimate of the work.) Thus, we are assuming that this board will take at least 4 moves to solve.

Here is another example

2	3	6		1	1	1
4	5	9		0	0	1
7	8	1		0	0	2

This board has a total distance of 6. However, since one rotation can affect the distance by 3, we divide the total distance by 3 so that we have an underestimate of the actual number of moves

required. Thus, we estimate this board requires at least two moves to solve. (Note, the board can be solved in exactly 2 moves.)

So (for this example) if we had already taken 5 moves to get to the current state, the expected cost for the state would be 7 (5 moves to get to the state and an estimated 2 more moves to finally win). Notice we are making a crude estimate of how good the state is. We prefer a state which (1) took few moves to reach and (2) is close to the goal state.

The priority of any board is the total expected moves.

We make a key observation: to solve the puzzle from a given state on the priority queue, the total number of moves required (including those already made) is at least our computed priority (expected moves)..

Consequently, as soon as we dequeue a board which has the goal state, we have not only discovered a sequence of moves from the initial board to the board associated with the state, but one that makes the fewest number of moves. **Notice that (in general) when we enqueue the state, we may not know we won't find another way which is better.** When we dequeue, we know everything else we will enqueue in the future will take at least as many moves to reach the goal. This is true because our distance is an UNDER-estimate of the number of moves required.

Output:

Since we want to compare this version with our brute force solution in program 1, you will need to have both methods working. (You may borrow a working version of program 1 from a friend if you don't have one.) Create a class "RotationGame" which allows you to call either "bruteForceSolve" or "aStarSolve". Each method should keep track of the total number of states that were pulled off the queue.

The user should be able to print the series of states that are **dequeued** from the priority queue, ending with the winning state.

Show the output for the following sample inputs:

Input 0

2 3 6

4 5 9

7 8 1

Input 1

5 6 4

8 2 9

1 7 3

Input 2

2 1 8

9 4 7

3 6 5

Input 3: Find an example for which your “more intelligent” method saves significant time over the brute force method of program 1. Enter the board for this example in your test code.

Sample Output

When using the A* search, allow the user to see each state as you pull it off the queue. This should be controlled by a Boolean variable “seeAll”. Be sure to print all the data in the state.

When using the brute force search, just show us the final sequence of moves and the total number of states dequeued.

The output below shows the cost so far in parens and expected cost of the board in square brackets.

The output might look like:

```
State 0 From State -1 (0):[2]
2 3 6
4 5 9
7 8 1
removed 0 states
State 3 From State 0 >0 (1):[2]
6 2 3
4 5 9
7 8 1
removed 1 states
State 10 From State 0 V2 (1):[2]
2 3 1
4 5 6
7 8 9
removed 2 states
State 27 From State 10 V2>0 (2):[2]
1 2 3
4 5 6
7 8 9
removed 3 states
```

Hints

Print output both to the console and to a file. Then, you can easily remove the printing to the console (to make it run faster and easier to read).

Getting the same code to work for ints and GameState forces us to be more methodical in our approach. For example:

The toString function looks something like:

```
string toString( AvlNode *t, string indent) const
{
    stringstream ss;
    if( t != nullptr )
    {
        ss << toString(t->right, indent + " ");
        ss << indent <<t->element << endl;
```

```

    ss << toString(t->left, indent + " ");
}
return ss.str();
}

```

This works fine if element is an int, but not so great if element is a GameState. How do you get this to work? Ints don't have an element component. GameStates don't know how to print themselves.

AHHHHH. You need to overload << for GameStates.

It is done like:

```

ostream& operator<<(ostream& ss, const GameState& gs) {
    ss << gs.toString() << endl;
    return ss;
}

```

Note however, that this cannot be a member function of GameState because the first parameter of << needs to be an ostream&. When you make something a member function of GameState, the first (understood) parameter is GameState. Right?

This is where a friend function comes in. (The function needs access to private data within GameState, as if it were a member.) Inside the GameState you declare:

```

friend ostream& operator<<( ostream& ss, const GameState& gs);

```

What to Turn in:

Submit a zip file containing your software project . The zip file should contain a readme file which tells how to run the program (what IDE and compiler you used).